Daylight - Data Engineering Task

Build a production-ready data pipeline that ingests audit logs from Cloudflare and makes them available for analytics and long-term storage.

- Continuously ingest audit logs from the Cloudflare Audit Logs API.
- Treat ingestion as a streaming data pipeline.
- Deliver the data to:
    - An analytical database for near-real-time querying
    - S3-compatible storage (e.g., MinIO) for long-term storage in an open, analytics-friendly format.

Notes

- Create a free Cloudflare account at cloudflare.com.
    - **API Documentation:** https://developers.cloudflare.com/api/resources/audit_logs/methods/list/
- You can choose any analytical database.
- The data must be queryable shortly after ingestion.
- For S3-compatible storage, you can use MinIO as a free, self-hosted alternative to Amazon S3.
- Storage should use an open format suitable for long-term storage and future analytical use.
    - The stored data should be easily loadable into analytical databases at a later time.
    - Use an open data format such as Iceberg, Parquet, ORC, or Delta Lake.

Expectations

- Submit a repository containing all relevant source code and configuration required to run the pipeline end-to-end.
- The solution should demonstrate production-grade quality — according to what you consider as production grade.